

Traditional programming

Inputs



Rules

1. Cut vegetables
2. Season chicken
3. Preheat oven
4. Cook chicken for 30-minutes
5. Add vegetables



Output



Starts with

Makes

Machine learning algorithm

Inputs



Output



Rules

1. Cut vegetables
2. Season chicken
3. Preheat oven
4. Cook chicken for 30-minutes
5. Add vegetables

Starts with

Figures out

(maybe not very simple...)

“If you can build a **simple rule-based** system that doesn't require machine learning, do that.”

— A wise software engineer... (actually rule 1 of Google's Machine Learning Handbook)



What deep learning is good for



- **Problems with long lists of rules**—when the traditional approach fails, machine learning/deep learning may help.
- **Continually changing environments**—deep learning can adapt ('learn') to new scenarios.
- **Discovering insights within large collections of data**—can you imagine trying to hand-craft rules for what 101 different kinds of food look like?

What deep learning is ^(typically) not good for

- **When you need explainability**—the patterns learned by a deep learning model are typically uninterpretable by a human.
- **When the traditional approach is a better option** — if you can accomplish what you need with a simple rule-based system.
- **When errors are unacceptable** — since the outputs of deep learning model aren't always predictable.
- **When you don't have much data** — deep learning models usually require a fairly large amount of data to produce great results.

(though we'll see how to get great results without huge amounts of data)

Machine Learning vs. Deep Learning

(common algorithms)

- Random forest
- Gradient boosted models
- Naive Bayes
- Nearest neighbour
- Support vector machine
- ...many more

(since the advent of deep learning these are often referred to as "shallow algorithms")

- Neural networks
- Fully connected neural network
- Convolutional neural network
- Recurrent neural network
- Transformer
- ...many more

What we're focused on building
(with PyTorch)

Structured data

*(depending how you represent your problem,
many algorithms can be used for both)*

Unstructured data

Neural Networks



(before data gets used with a neural network, it needs to be turned into numbers)

Daniel Bourke @mrdbourke · Nov 1
"How do I learn #machinelearning?"

What you want to hear:

1. Learn Python
2. Learn Math/Stats/Probability
3. Learn software engineering
4. Build

What you need to do:

1. Google it
2. Go down the rabbit hole
3. Resurface in 6-9 months and reassess

See you on the other side.



Inputs

Numerical encoding

$\begin{bmatrix} 116, & 78, & 15 \\ 117, & 43, & 96 \\ 125, & 87, & 23 \end{bmatrix},$
...

Each of these nodes is called a "hidden unit" or "neuron".



(choose the appropriate neural network for your problem)

Learns representation (patterns/features/weights)

$\begin{bmatrix} 0.983, & 0.004, & 0.013 \\ 0.110, & 0.889, & 0.001 \\ 0.023, & 0.027, & 0.985 \end{bmatrix},$
...

Representation outputs

(a human can understand these)

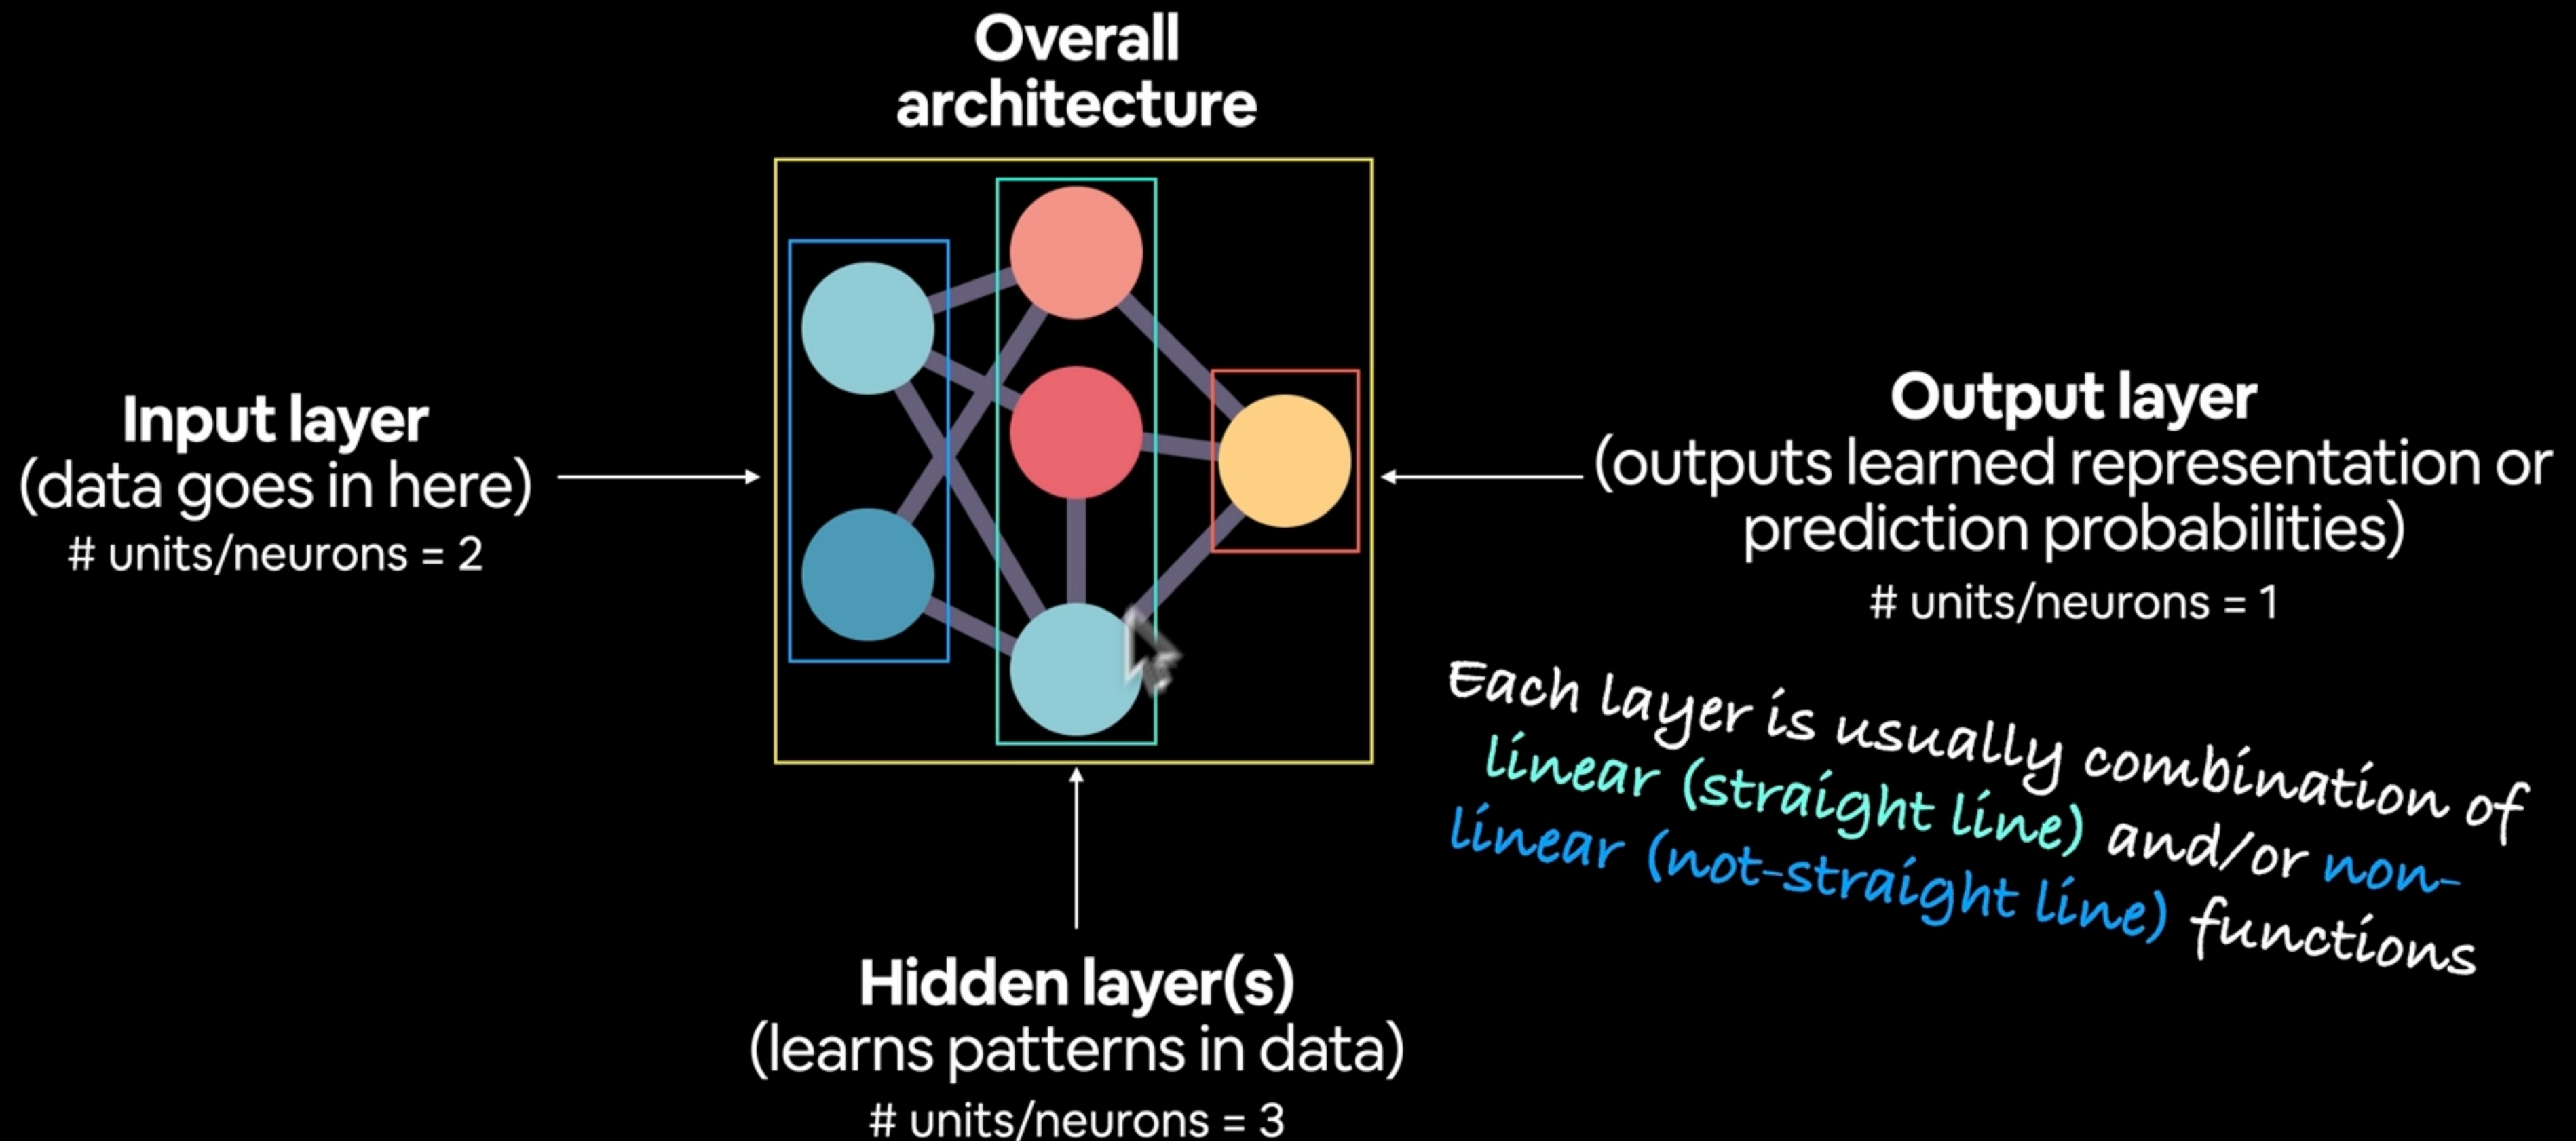
Ramen, Spaghetti

Not a diaster

"Hey Siri, what's the weather today?"

Outputs

Anatomy of Neural Networks



Note: “patterns” is an arbitrary term, you’ll often hear “embedding”, “weights”, “feature representation”, “feature vectors” all referring to similar things.


Types of Learning



**Supervised
Learning**



**Unsupervised &
Self-supervised
Learning**



**Transfer
Learning**

We'll be writing code to do these,
but the style of code can be adopted across learning paradigms.